

**function** TREE-SEARCH(*problem*) **returns** a solution, or failure  
    initialize the frontier using the initial state of *problem*  
    **loop do**  
        **if** the frontier is empty **then return** failure  
        choose a leaf node and remove it from the frontier  
        **if** the node contains a goal state **then return** the corresponding solution  
        expand the chosen node, adding the resulting nodes to the frontier

**function** GRAPH-SEARCH(*problem*) **returns** a solution, or failure  
    initialize the frontier using the initial state of *problem*  
    **initialize the explored set to be empty**  
    **loop do**  
        **if** the frontier is empty **then return** failure  
        choose a leaf node and remove it from the frontier  
        **if** the node contains a goal state **then return** the corresponding solution  
        **add the node to the explored set**  
        expand the chosen node, adding the resulting nodes to the frontier  
        **only if not in the frontier or explored set**

**function** UNIFORM-COST-SEARCH(*problem*) **returns** a solution, or failure  
    *node*  $\leftarrow$  a node with STATE = *problem*.INITIAL-STATE, PATH-COST = 0  
    *frontier*  $\leftarrow$  a priority queue ordered by PATH-COST, with *node* as the only element  
    *explored*  $\leftarrow$  an empty set  
    **loop do**  
        **if** EMPTY?(*frontier*) **then return** failure  
        *node*  $\leftarrow$  POP(*frontier*)   /\* chooses the lowest-cost node in *frontier* \*/  
        **if** *problem*.GOAL-TEST(*node*.STATE) **then return** SOLUTION(*node*)  
        add *node*.STATE to *explored*  
        **for each action in** *problem*.ACTIONS(*node*.STATE) **do**  
            *child*  $\leftarrow$  CHILD-NODE(*problem*, *node*, *action*)  
            **if** *child*.STATE is not in *explored* or *frontier* **then**  
                *frontier*  $\leftarrow$  INSERT(*child*, *frontier*)  
            **else if** *child*.STATE is in *frontier* with higher PATH-COST **then**  
                replace that *frontier* node with *child*

**function** DEPTH-LIMITED-SEARCH(*problem*, *limit*) **returns** a solution, or failure/cutoff  
    **return** RECURSIVE-DLS(MAKE-NODE(*problem*.INITIAL-STATE), *problem*, *limit*)

**function** RECURSIVE-DLS(*node*, *problem*, *limit*) **returns** a solution, or failure/cutoff  
    **if** *problem*.GOAL-TEST(*node*.STATE) **then return** SOLUTION(*node*)  
    **else if** *limit* = 0 **then return** *cutoff*  
    **else**  
        *cutoff\_occurred?*  $\leftarrow$  false  
        **for each action in** *problem*.ACTIONS(*node*.STATE) **do**  
            *child*  $\leftarrow$  CHILD-NODE(*problem*, *node*, *action*)  
            *result*  $\leftarrow$  RECURSIVE-DLS(*child*, *problem*, *limit* - 1)  
            **if** *result* = *cutoff* **then** *cutoff\_occurred?*  $\leftarrow$  true  
            **else if** *result*  $\neq$  failure **then return** *result*  
        **if** *cutoff\_occurred?* **then return** *cutoff* **else return** failure

**function** ITERATIVE-DEEPENING-SEARCH(*problem*) **returns** a solution, or failure  
    **for** *depth* = 0 to  $\infty$  **do**  
        *result*  $\leftarrow$  DEPTH-LIMITED-SEARCH(*problem*, *depth*)  
        **if** *result*  $\neq$  cutoff **then return** *result*

**function** MINIMAX-DECISION(*state*) **returns** an action  
    **return** arg max<sub>*a*  $\in$  ACTIONS(*s*)</sub> MIN-VALUE(RESET(*state*, *a*))

**function** MAX-VALUE(*state*) **returns** a utility value  
    **if** TERMINAL-TEST(*state*) **then return** UTILITY(*state*)  
    *v*  $\leftarrow -\infty$   
    **for each** *a* **in** ACTIONS(*state*) **do**  
        *v*  $\leftarrow$  MAX(*v*, MIN-VALUE(RESET(*s*, *a*)))  
    **return** *v*

**function** MIN-VALUE(*state*) **returns** a utility value  
    **if** TERMINAL-TEST(*state*) **then return** UTILITY(*state*)  
    *v*  $\leftarrow \infty$   
    **for each** *a* **in** ACTIONS(*state*) **do**  
        *v*  $\leftarrow$  MIN(*v*, MAX-VALUE(RESET(*s*, *a*)))  
    **return** *v*

**function** ALPHA-BETA-SEARCH(*state*) **returns** an action  
    *v*  $\leftarrow$  MAX-VALUE(*state*,  $-\infty$ ,  $+\infty$ )  
    **return** the action in ACTIONS(*state*) with value *v*

**function** MAX-VALUE(*state*,  $\alpha$ ,  $\beta$ ) **returns** a utility value  
    **if** TERMINAL-TEST(*state*) **then return** UTILITY(*state*)  
    *v*  $\leftarrow -\infty$   
    **for each** *a* **in** ACTIONS(*state*) **do**  
        *v*  $\leftarrow$  MAX(*v*, MIN-VALUE(RESET(*s*, *a*),  $\alpha$ ,  $\beta$ ))  
        **if** *v*  $\geq \beta$  **then return** *v*  
         $\alpha \leftarrow$  MAX( $\alpha$ , *v*)  
    **return** *v*

**function** MIN-VALUE(*state*,  $\alpha$ ,  $\beta$ ) **returns** a utility value  
    **if** TERMINAL-TEST(*state*) **then return** UTILITY(*state*)  
    *v*  $\leftarrow +\infty$   
    **for each** *a* **in** ACTIONS(*state*) **do**  
        *v*  $\leftarrow$  MIN(*v*, MAX-VALUE(RESET(*s*, *a*),  $\alpha$ ,  $\beta$ ))  
        **if** *v*  $\leq \alpha$  **then return** *v*  
         $\beta \leftarrow$  MIN( $\beta$ , *v*)  
    **return** *v*

Solving Problems by Searching

- Define search problem
- State space (should be partial-valid)
  - Initial state
  - Goal state(s) / goal test
  - Actions(s) returns actions available to agent at s
  - Transition model Result(s, a) returns next state
  - Action cost function Action-Cost(s, a, s')
  - Path cost

Formulating problems: modelling the search problem through abstraction

Note: assume goal exists at finite depth

- Tree and Graph Search Algo
- Start at initial state, keep searching till reaching a goal state
  - Frontier: nodes that we have seen but haven't explored yet. At initialisation, frontier is just the source.
  - At each iteration, choose a node from frontier, explore it, and add its neighbours to frontier
  - Graph search: a node that's been explored once will not be revisited.
  - A state: represents a physical config
  - A node: a data structure constituting part of search tree. It includes state, parent node, action, and path cost g(n).
  - 2 diff nodes can contain the same world state.

- Search problem params
- b: branching factor
  - d: depth of shallowest goal node
  - m: maximum depth of search tree (may be inf)

Uninformed Search

- Breadth-First Search (BFS)
- Expand shallowest unexpanded node
  - Frontier is FIFO queue
  - Goal test is applied when pushing nodes to frontier rather than during expansion
  - # of nodes:  $O(b) + O(b^2) + \dots + O(b^d)$

- Uniform-Cost Search (UCS)
- Expand least-path-cost unexpanded node
  - Frontier is PQ ordered by path cost
  - Equivalent to BFS if all step costs are equal

- Depth-First Search (DFS)
- Expand least-deepest unexpanded node
  - Frontier is LIFO stack

- Depth-Limited Search (DLS)
- Run DFS with depth limit l

- Iterative Deepening Search (IDS)
- Perform DLSs with increasing depth limit until goal node is found
  - Better if state space is large and depth of solution is unknown
  - # nodes:  $(d + 1)O(b^0) + dO(b^1) + (d - 1)O(b^2) + \dots + 2O(b^{d-1}) + O(b^d)$

Property	BFS	UCS	DFS	DLS	IDS
Complete	Yes*	Yes*	No	No	Yes*
Optimal	No	Yes	No	No	No
Time	$O(b^d)$	$O(b^{1+\lceil \frac{c^*}{\epsilon} \rceil})$	$O(b^m)$	$O(b^l)$	$O(b^d)$
Space	$O(b^d)$	$O(b^{1+\lceil \frac{c^*}{\epsilon} \rceil})$	$O(bm)$	$O(bl)$	$O(bd)$
1. BFS and IDS are complete if b is finite.					
2. UCS is complete if b is finite and step cost $\geq \epsilon$					
3. BFS and IDS are optimal if step costs are identical.					

- Proof of UCS' optimality
- Let  $c(n)$  be the cost of the path to node  $n$ .  
If  $n_2$  is expanded after  $n_1$ , then  $c(n_1) \leq c(n_2)$
- Case 1:  $n_2$  is on the frontier when  $n_1$  is expanded
  - Case 2:  $n_2$  was added to the frontier when  $n_1$  was expanded
- When  $n$  is expanded, every path with cost  $< c(n)$  has already been expanded.
- Let  $S_0, n_0, n_1, \dots, n_k$  be a path with cost  $< c(n)$ . Let  $n_i$  be the last node on this path that has been expanded.
  - $n_i + 1$  is still on the frontier. And  $c(n_i + 1) < c(n)$ .
  - UCS would have expanded  $n_i + 1$ , not  $n$ . So every node on this path must already be expanded.
- The first time UCS expands a state, it has found the minimal cost path to it
- No cheaper path exists, else that path would have been expanded before.
  - No cheaper path will be discovered later, as all those paths must be at least as expensive.

BFS	<ul style="list-style-type: none"><li>- Goal node is near root</li><li>- Tree is deep but goals are rare</li></ul>
DFS	<ul style="list-style-type: none"><li>- Goal node is very deep or all goal nodes are at the same depth</li><li>- Better space complexity than BFS</li></ul>
UCS	<ul style="list-style-type: none"><li>- If cost is known/non-uniform and optimality is a requirement</li><li>- Equivalent to BFS if step costs are uniform</li></ul>
LDFS	<ul style="list-style-type: none"><li>- If we know at what depth the goal node is</li></ul>
IDS	<ul style="list-style-type: none"><li>- Like a fusion of BFS and DFS</li><li>- Some overhead (1/(b-1))</li></ul>

**Informed Search**

**Best-First Search**

- Use evaluation function  $f(n)$  as a cost estimate
- Frontier: PQ ordered by non-decreasing cost  $f$

**Greedy Best-First Search**

- $f(n) = h(n)$
- $h(n)$ : heuristic function, which estimates the cheapest path cost from  $n$  to goal
- Greedy best-first search expands the node that appears closest to goal
- Completeness: if  $b$  is finite, tree-based variant is incomplete, while graph-based variant is complete
- Not optimal
- Time & space  $O(b^m)$

**A\* Search**

- $f(n) = g(n) + h(n)$
- $g(n)$ : cost of reaching  $n$  from start node
- Complete if finite # of nodes and  $f(n) \leq f(G)$
- Optimality depends on heuristics
- Time  $O(b^{h^*(s_0) - h(s_0)})$
- Space  $O(b^m)$

**Admissible Heuristic**

- Admissible heuristic:  $\forall n, h(n) \leq h^*(n)$ , i.e. never overestimates cost to reach goal
- If  $h(n)$  is admissible, then A\* using Tree-Search is optimal

Proof.

- If A\* using admissible heuristic returns suboptimal goal  $t$ , then there exists a node  $n$  in frontier, on optimal path but not expanded.
- $f(t) = g(t) > g^*(t) = f^*(t) = g(n) + h^*(n) \geq g(n) + h(n) = f(n)$
- Admissible heuristic doesn't guarantee optimality for Graph-search. Graph-search discards new paths to a repeated state. If the heuristic is not consistent, the optimal path might be discarded

Dominance: If  $h_2(n) \geq h_1(n)$  for all  $n$ , then  $h_2$  dominates  $h_1$ . it follows that  $h_2$  incurs lower search cost than  $h_1$ .

Deriving Admissible Heuristics:

- Relaxed problem: one with fewer restrictions on actions
- The cost of an optimal solution to a relaxed problem is an admissible heuristic for the original proble

**Consistent Heuristic**

- Consistent heuristic: for every node  $n$  and successor  $n'$  of  $n$  generated by  $a$ ,  $h(n) \leq d(n, n') + h(n')$
  - Equivalently,  $f(n)$  is non-decreasing along any path.  $f(n') = g(n') + h(n') = g(n) + d(n, n') + h(n') \geq g(n) + h(n) = f(n)$
  - Consistency implies admissibility (proof by induction)
  - If  $h(n)$  is consistent, then A\* using Graph-Search is optimal
- Proof. When A\* selects a node  $n$  for expansion, the shortest path to  $n$  has been found
- If A\* returns a suboptimal path to  $n$ , there exists a node  $m$  in frontier, on optimal path but not expanded.
  - However, A\* with consistent heuristic explores nodes in a non-decreasing order of  $f$  value, hence  $m$  should have been explored before  $n$ .

**Local Search**

- The path to goal is irrelevant; the goal state itself is the solution
- State space: set of complete configs
- Find final configs satisfying constraints
- Local search algo: maintain single current best state and try to improve it
- Advantages: very little/constant memory, and can find reasonable solutions in large state space

**Hill-climbing search**

- if highest-valued successor if better than current, update current.
- Always terminate with a solution
- Problem: can get stuck in local maximum
- Non-guaranteed fixes: sideways moves, random restarts

**Adversarial Search aka Games**

**Game: Problem Formulation**

- Initial state
- States
- Players: Player(s) defines which player has the move in state  $s$
- Actions: Actions(s) returns the set of legal moves in  $s$
- Transition model: Result( $s, a$ )
- Terminal test Terminal( $s$ ) == true iff game end
- Utility function Utility( $s, p$ ): final numeric value for a game that ends in terminal state  $s$  for player  $p$

**Winning Strategy**

- Let  $V_{\max}$  be the set of nodes controlled by the MAX player and  $V_{\min}$  be the set of nodes controlled by the MIN player.
- A *strategy* for the MAX player is a mapping  $s_1 : V_{\max} \rightarrow V$ ; similarly, a strategy for the MIN player is a mapping  $s_2 : V_{\min} \rightarrow V$ .
- A strategy  $s_1^*$  for player 1 is called winning if for any strategy  $s_2$  by player 2, the game ends with player 1 as the winner.
- The leaves of the minimax tree are *payoff nodes*. There is a payoff  $a(v) \in \mathbb{R}$  associated with each payoff node  $v$ . More formally, the utility of the MAX player from  $v$  is  $u_{\max}(v) = a(v)$  and the utility of the MIN player is  $u_{\min}(v) = -a(v)$ . The utility of a player from a pair of strategies  $s_1 \in S_1, s_2 \in S_2$  is simply the utility they receive by the leaf node reached when the strategy pair  $(s_1, s_2)$  is played.

**Optimal Strategy at Node - Minimax**

- MAX chooses move to maximize the minimum payoff
- MIN chooses move to minimize the maximum payoff

$$\begin{aligned}
 & \text{Minimax}(s) \\
 &= \begin{cases} \text{Utility}(s) & \text{if TerminalTest}(s) \\ \max_{a \in \text{Actions}(s)} \text{Minimax}(\text{Result}(s, a)) & \text{if Player}(s) = \text{MAX} \\ \min_{a \in \text{Actions}(s)} \text{Minimax}(\text{Result}(s, a)) & \text{if Player}(s) = \text{MIN} \end{cases}
 \end{aligned}$$

- Complete if game tree is finite
- Optimal
- Time  $O(b^m)$
- Space  $O(bm)$
- Returns a SPNE: best action at every choice node

Proof: by induction

- Assume MINIMAX computes SPNE for all subtrees at height  $h-1$ . WLOG, consider node  $v$  at height  $h$  and assume this is a MAX node.
- Let  $s_1^*$  be the strategy outputted by MINIMAX, and  $s_1$  another strategy by player 1
- Suppose that  $v_1^*, v_1$  are the nodes chosen by  $s_1^*$  and  $s_1$
- $u_1(v, s_1^*, s_2^*) = u_1(v_1^*, s_1^*, s_2^*) \geq u_1(v_1, s_1^*, s_2^*) \geq u_1(v_1, s_1, s_2^*) = u_1(v, s_1, s_2^*)$

**Alpha-beta pruning**

- Maintain a lower bound  $\alpha$  and upper bound  $\beta$  of the values of MAX's and MIN's nodes seen thus far
- MAX node  $n$ :  $\alpha(n)$  = highest observed value found on path from  $n$ ; initially  $\alpha(n) = -\infty$
- MIN node  $n$ :  $\beta(n)$  = lowest observed value found on path from  $n$ ; initially  $\beta(n) = +\infty$
- Given a MIN node  $n$ , stop searching below  $n$  if there is some MAX ancestor  $i$  of  $n$  with  $\alpha(i) \geq \beta(n)$
- Given a MAX node  $n$ , stop searching below  $n$  if there is some MIN ancestor  $i$  of  $n$  with  $\beta(i) \leq \alpha(n)$
- Pruning never affects the final outcome i.e. it leaves at least one strategy played in a Nash Equilibrium; however, alpha-beta pruning cannot be used to find SPNE.
- Time:  $O\left(b^{\frac{m}{2}}\right)$  for perfect ordering,  $O\left(b^{\frac{3m}{4}}\right)$  for random ordering when  $b < 1000$

**Evaluation function and cut-off test**

- Evaluation function: estimated expected utility of state
- Cut-off test: depth limit
- Heuristic minimax value: run minimax until depth  $d$ , then start evaluation function to choose nodes

$$\begin{aligned}
 & \text{H-MINIMAX}(s, d) = \\
 & \begin{cases} \text{EVAL}(s) & \text{if CUTOFF-TEST}(s, d) \\ \max_{a \in \text{ACTIONS}(s)} \text{H-MINIMAX}(\text{RESULT}(s, a), d + 1) & \text{if PLAYER}(s) = \text{MAX} \\ \min_{a \in \text{ACTIONS}(s)} \text{H-MINIMAX}(\text{RESULT}(s, a), d + 1) & \text{if PLAYER}(s) = \text{MIN} \end{cases}
 \end{aligned}$$