

NATIONAL UNIVERSITY OF SINGAPORE

SCHOOL OF COMPUTING

MIDTERM ASSESSMENT FOR
Semester III AY2019/20

CS3243: INTRODUCTION TO ARTIFICIAL INTELLIGENCE

June 4, 2020

Time Allowed: 2 Hours

Part I: Summary of Course Material

The following summary contains key parts from the course lecture notes. It **does not** offer a complete coverage of course materials. You may use any of the claims shown here without proving them, unless specified explicitly in the question.

I.1 Search Problems**I.1.1 Uninformed Search**

We are interested in finding a solution to a fully observable, deterministic, discrete problem. A search problem is given by a set of *states*, where a *transition function* $T(s, a, s')$ states that taking action a in state s will result in transitioning to state s' . There is a cost to this defined as $c(s, a, s')$, assumed to be non-negative. We are also given an initial state (assumed to be in our frontier upon initialization). The *frontier* is the set of nodes in a queue that have not yet been explored, but will be explored given the order of the queue. We also have a *goal test*, which for a given state s outputs “yes” if s is a goal state (there can be more than one).

```

function TREE-SEARCH(problem) returns a solution, or failure
  initialize the frontier using the initial state of problem
  loop do
    if the frontier is empty then return failure
    choose a leaf node and remove it from the frontier
    if the node contains a goal state then return the corresponding solution
    expand the chosen node, adding the resulting nodes to the frontier

```

Figure I.1: The tree search algorithm

```

function GRAPH-SEARCH(problem) returns a solution, or failure
  initialize the frontier using the initial state of problem
  initialize the explored set to be empty
  loop do
    if the frontier is empty then return failure
    choose a leaf node and remove it from the frontier
    if the node contains a goal state then return the corresponding solution
    add the node to the explored set
    expand the chosen node, adding the resulting nodes to the frontier
    only if not in the frontier or explored set

```

Figure I.2: The graph search algorithm

We have discussed two search variants in class, tree search (Figure I.1) and graph search (Figure I.2). They differ in the fact that under graph search we do not explore nodes that we have seen before.

The main thing that differentiates search algorithms is the order in which we explore the frontier. In breadth-first search we explore the shallowest nodes first; in depth-first search we explore the deepest nodes first; in uniform-cost search (Figure I.3) we explore nodes in order of cost.

```

function UNIFORM-COST-SEARCH(problem) returns a solution, or failure
  node  $\leftarrow$  a node with STATE = problem.INITIAL-STATE, PATH-COST = 0
  frontier  $\leftarrow$  a priority queue ordered by PATH-COST, with node as the only element
  explored  $\leftarrow$  an empty set
  loop do
    if EMPTY?(frontier) then return failure
    node  $\leftarrow$  POP(frontier) /* chooses the lowest-cost node in frontier */
    if problem.GOAL-TEST(node.STATE) then return SOLUTION(node)
    add node.STATE to explored
    for each action in problem.ACTIONS(node.STATE) do
      child  $\leftarrow$  CHILD-NODE(problem, node, action)
      if child.STATE is not in explored or frontier then
        frontier  $\leftarrow$  INSERT(child, frontier)
      else if child.STATE is in frontier with higher PATH-COST then
        replace that frontier node with child

```

Figure I.3: The uniform cost search algorithm

We have also studied variants where we only run DFS to a certain depth (Figure I.4) and where we iteratively deepen our search depth (Figure I.5). Table 1 summarizes the various search algorithms' properties.

```

function DEPTH-LIMITED-SEARCH(problem, limit) returns a solution, or failure/cutoff
  return RECURSIVE-DLS(MAKE-NODE(problem.INITIAL-STATE), problem, limit)

function RECURSIVE-DLS(node, problem, limit) returns a solution, or failure/cutoff
  if problem.GOAL-TEST(node.STATE) then return SOLUTION(node)
  else if limit = 0 then return cutoff
  else
    cutoff_occurred?  $\leftarrow$  false
    for each action in problem.ACTIONS(node.STATE) do
      child  $\leftarrow$  CHILD-NODE(problem, node, action)
      result  $\leftarrow$  RECURSIVE-DLS(child, problem, limit - 1)
      if result = cutoff then cutoff_occurred?  $\leftarrow$  true
      else if result  $\neq$  failure then return result
    if cutoff_occurred? then return cutoff else return failure

```

Figure I.4: The depth-limited search algorithm

```

function ITERATIVE-DEEPENING-SEARCH(problem) returns a solution, or failure
  for depth = 0 to  $\infty$  do
    result  $\leftarrow$  DEPTH-LIMITED-SEARCH(problem, depth)
    if result  $\neq$  cutoff then return result

```

Figure I.5: The iterative deepening search algorithm

We also noted that no deterministic search algorithm can do well in general; in the worst case, they will all search through the entire search space.

| Property | BFS | UCS | DFS | DLS | IDS |
|--|--------------------|---|--------------------|-----------------------|--------------------|
| Complete | Yes* | Yes* | No | No | Yes* |
| Optimal | No | Yes | No | No | No |
| Time | $\mathcal{O}(b^d)$ | $\mathcal{O}(b^{1+\lfloor \frac{C^*}{\epsilon} \rfloor})$ | $\mathcal{O}(b^m)$ | $\mathcal{O}(b^\ell)$ | $\mathcal{O}(b^d)$ |
| Space | $\mathcal{O}(b^d)$ | $\mathcal{O}(b^{1+\lfloor \frac{C^*}{\epsilon} \rfloor})$ | $\mathcal{O}(bm)$ | $\mathcal{O}(b\ell)$ | $\mathcal{O}(bd)$ |
| * Under certain assumptions made on the search space | | | | | |

Table 1: Summary of search algorithms' properties

I.1.2 Informed Search

Informed search uses additional information about the underlying search problem in order to narrow down the search scope. We have mostly discussed the A^* algorithm, where the priority queue holding the frontier is ordered by $g(v) + h(v)$, where $g(v)$ is the distance of a node from the source, and $h(v)$ is a *heuristic estimate* of the distance of the node from a goal node. There are two key types of heuristic functions

Definition I.1. A heuristic h is *admissible* if it never overestimates the distance of a node from the nearest goal; i.e.

$$\forall v : h(v) \leq h^*(v),$$

where $h^*(v)$ is the *optimal heuristic*, i.e. the true distance of a node from the nearest goal.

Definition I.2. A heuristic h is *consistent* if it satisfies the triangle inequality; i.e.

$$\forall v, v' : h(v) \leq h(v') + c(v, a, v')$$

where $c(v, a, v')$ is the cost of transitioning from v to v' via action a .

We have shown that A^* with tree search is optimal when the heuristic is admissible, and is optimal with graph search when the heuristic is consistent. We also showed that consistency implies admissibility but not the other way around, and that running A^* with an admissible inconsistent heuristic with graph search may lead to a sub-optimal goal.

I.2 Adversarial Search

An extensive form game is defined by V a set of nodes and E a set of directed edges, defining a tree. The root of the tree is the node r . Let V_{\max} be the set of nodes controlled by the MAX player and V_{\min} be the set of nodes controlled by the MIN player. We often refer to the MAX player as player 1, and to the MIN player as player 2. A *strategy* for the MAX player is a mapping $s_1 : V_{\max} \rightarrow V$; similarly, a strategy for the MIN player is a mapping $s_2 : V_{\min} \rightarrow V$. In both cases, $s_i(v) \in \text{chld}(v)$ is the choice of child node that will be taken at node v . We let $\mathcal{S}_1, \mathcal{S}_2$ be the set of strategies for the MAX and MIN player, respectively.

The leaves of the minimax tree are *payoff nodes*. There is a payoff $a(v) \in \mathbb{R}$ associated with each payoff node v . More formally, the utility of the MAX player from v is $u_{\max}(v) = a(v)$ and the utility of the MIN player is $u_{\min}(v) = -a(v)$. The utility of a player from a pair of strategies $s_1 \in \mathcal{S}_1, s_2 \in \mathcal{S}_2$ is simply the utility they receive by the leaf node reached when the strategy pair (s_1, s_2) is played.

Definition I.3 (Nash Equilibrium). A pair of strategies $s_1^* \in \mathcal{S}_1, s_2^* \in \mathcal{S}_2$ for the MAX player and the MIN player, respectively, is a *Nash equilibrium* if no player can get a strictly higher utility by switching their strategy. In other words:

$$\begin{aligned} \forall s \in \mathcal{S}_1 : u_1(s_1^*, s_2^*) &\geq u_1(s, s_2^*); \\ \forall s' \in \mathcal{S}_2 : u_2(s_1^*, s_2^*) &\geq u_2(s_1^*, s') \end{aligned}$$

Definition I.4 (Subgame). Given an extensive form game $\langle V, E, r, V_{\max}, V_{\min}, \vec{a} \rangle$, a subgame is a subtree of the original game, defined by some arbitrary node v set to be the root node r , and all of its descendants (i.e. its children, its children's children etc.), denoted by $\text{desc}(v)$. Terminal node payoffs the same as in the original extensive form game, and players still control the same nodes as in the original game.

Definition I.5 (Subgame-Perfect Nash Equilibrium (SPNE)). A pair of strategies $s_1^* \in \mathcal{S}_1, s_2^* \in \mathcal{S}_2$ is a subgame-perfect Nash equilibrium if it is a Nash equilibrium for any subtree of the original game tree.

Figure I.6 describes the minimax algorithm, which computes SPNE strategies for the MIN and MAX players, as we have shown in class. We discussed the α - β pruning algorithm as a method of removing subtrees that need not be explored (Figure I.7).

| |
|---|
| function MINIMAX-DECISION(<i>state</i>) returns an action return $\arg \max_{a \in \text{ACTIONS}(s)} \text{MIN-VALUE}(\text{RESULT}(\text{state}, a))$ |
| function MAX-VALUE(<i>state</i>) returns a utility value if TERMINAL-TEST(<i>state</i>) then return UTILITY(<i>state</i>) $v \leftarrow -\infty$ for each <i>a</i> in ACTIONS(<i>state</i>) do $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a)))$ return <i>v</i> |
| function MIN-VALUE(<i>state</i>) returns a utility value if TERMINAL-TEST(<i>state</i>) then return UTILITY(<i>state</i>) $v \leftarrow \infty$ for each <i>a</i> in ACTIONS(<i>state</i>) do $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a)))$ return <i>v</i> |

Figure I.6: The minimax algorithm (note a typo from the AIMA book - *s* should be *state*).

| |
|--|
| function ALPHA-BETA-SEARCH(<i>state</i>) returns an action $v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$ return the action in ACTIONS(<i>state</i>) with value <i>v</i> |
| function MAX-VALUE(<i>state</i> , α , β) returns a utility value if TERMINAL-TEST(<i>state</i>) then return UTILITY(<i>state</i>) $v \leftarrow -\infty$ for each <i>a</i> in ACTIONS(<i>state</i>) do $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$ if $v \geq \beta$ then return <i>v</i> $\alpha \leftarrow \text{MAX}(\alpha, v)$ return <i>v</i> |
| function MIN-VALUE(<i>state</i> , α , β) returns a utility value if TERMINAL-TEST(<i>state</i>) then return UTILITY(<i>state</i>) $v \leftarrow +\infty$ for each <i>a</i> in ACTIONS(<i>state</i>) do $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$ if $v \leq \alpha$ then return <i>v</i> $\beta \leftarrow \text{MIN}(\beta, v)$ return <i>v</i> |

Figure I.7: The α - β pruning algorithm.